

MSSE Capstone Project

Final Production Validation, Evidence Closeout & Presentation Readiness Meeting Minutes

TrustAI Marketplace

Date	September 3, 2026
Time	Approximately 10:08 AM PDT / 7:08 PM CAT
Duration	Approximately 45 minutes
Location	Virtual team meeting
Meeting Type	Final production validation, Capstone evidence closeout, presentation structure, rehearsal planning and submission-readiness alignment
Prepared For	TrustAI Marketplace MSSE Capstone team records

1. Attendance and Participation

Four team members participated in the substantive meeting discussion: Ahmed Al-Mandalawi, Ranga Nyamadzawo, Mulima Chibuye and Adrian Muchatibaya. Samar El Ghandour was not present for the substantive discussion and the team agreed to pass the rehearsal schedule and follow-up information to her afterward.

Team member	Participation	Role / participation reflected in this meeting
Ahmed Al-Mandalawi	Present	Product Owner / AI Analysis Lead. Presented the production-validation update, Capstone evidence review, model-evaluation work and presentation-preparation priorities.
Ranga Nyamadzawo	Present	Backend/platform and architecture contributor. Helped frame feature-freeze discipline, issue triage, presentation flow and the need to focus on assessed engineering evidence.
Mulima Chibuye	Present	Scrum Master / Project Manager. Participated in Trello/access, submission, timing, rehearsal and presentation-structure discussions.
Adrian Muchatibaya	Present	Frontend lead. Participated in scheduling, frontend/scope discussion and presentation planning.

Team member	Participation	Role / participation reflected in this meeting
Samar El Ghandour	Not present for substantive discussion	QA / DevOps lead. The team planned to pass the agreed rehearsal schedule and follow-up information to her.

2. Meeting Objective

The meeting was held after the major implementation and release work had effectively reached the submission stage. The group's purpose was to decide what still needed attention, verify that the final production application was in a good enough state to stop adding scope, organize the evidence required for the Capstone submission, and turn the August presentation plan into a concrete rehearsal and recording schedule.

- Review the latest production state and any remaining presentation-blocking concerns.
- Confirm that new feature development should stop unless a serious defect is discovered.
- Review the Capstone Handbook and rubric against the repository, documentation, Trello board, meeting records and sprint evidence.
- Explain the model-evaluation work and decide how it can support the project's engineering story.
- Locate missing or scattered sprint-demo and agreement artifacts before submission.
- Clean up the Trello board so ownership, completion, deferred work and future work match what was actually delivered.
- Agree on one coordinated 15–20 minute group presentation, including IDs, timing, handoffs and rehearsal.
- Set a concrete next-session time and assign immediate preparation work.

3. Project Position at the Start of the Meeting

The tone at the start of the call was noticeably different from the August 26 meeting. The team was no longer discussing whether the project was close to release; it was treating the product as essentially complete for Capstone purposes and concentrating on submission quality. Ahmed reported that he had tested the latest production version the previous night, including both the normal listing-analysis path and Visual Inspection. Other members had not all retested the latest merged version yet.

This continued the final-stage direction agreed on August 26: protect the release, avoid unnecessary new scope, keep legitimate limitations visible, and spend the remaining time on validation, documentation, task-board accuracy and the recorded demonstration.

Working position reached

The practical consensus was that further development should be limited to serious or presentation-blocking defects. Everything else should be handled through evidence closeout, documentation, issue tracking or future-work treatment.

4. Final Production Validation Update

4.1 Live production testing

Ahmed reported that he had exercised the final production build extensively the previous night and had tested both the primary text-analysis experience and the Visual Inspection path live. His view was that the application was now in a strong enough state for the team to move away from feature development and toward final submission work.

The group was asked to test the application independently as well, especially because some members had not yet reviewed the newest merged changes. The purpose of the remaining testing was not to restart broad development; it was to surface only defects that could materially affect the demonstration or the credibility of the final deliverable.

4.2 Production model and provider-agnostic architecture

A question was raised about whether the production text-analysis path was now tied to one specific model and how that would sit alongside the provider-agnostic architecture discussed throughout the project. The meeting did not establish a new architecture decision. The underlying concern was how to explain the release model clearly while preserving the distinction between a specific production choice and the broader provider abstraction built into the application.

The group treated this as something to explain clearly in the presentation rather than as a reason to reopen the architecture. A production choice can be specific while the software design remains provider-independent.

5. Feature Freeze and Remaining Development

Ahmed stated that, from his side, anything he would still like to add after this point belonged after the presentation unless the team found a serious defect. Ranga supported the same general direction: with the product in a good state, the team would gain more by getting the submission evidence and presentation in order than by starting another large piece of implementation.

- Do not begin discretionary feature work during the submission window.
- Continue testing the live application so serious defects are still found before the recording.
- If a newly discovered problem is small and clearly worth fixing, raise it and review the effort before changing the release.
- If a problem requires substantial work and does not block the agreed user stories or demonstration, keep it visible as future work or a known limitation.
- Do not artificially make the project look complete by erasing legitimate engineering follow-up.

This is consistent with the earlier August 26 agreement that feature freeze means freezing new scope, not pretending that software has no remaining defects or possible improvements.

6. Capstone Handbook and Rubric Evidence Closeout

Ahmed explained that he had gone back through the Capstone Handbook and rubric and had started comparing the repository and documentation against the grading expectations. His conclusion was that most of the remaining work was no longer application development. The remaining risk was whether the team's engineering work would be easy for the grader to find, understand and verify.

- Make the final documentation accurate and consistent with the released application.
- Bring the Trello board into line with the work actually completed.
- Organize the formal meeting minutes and sprint records so the project chronology is clear.
- Locate and index the sprint-demo recordings that genuinely exist.

- Complete the detailed design and testing document required by the Handbook.
- Put most of the remaining team effort into a strong, rehearsed demonstration.

The Handbook cross-check confirms why these items matter: the repository must link the deployed application and an accessible agile task board, include a detailed design-and-testing document, and be shared with the designated grader account. The final demonstration must be a single 15–20 minute recording in which all group members speak, are visible on camera, and present government-issued identification.

7. Model Evaluation as Engineering Evidence

7.1 Why the evaluation work matters

Ahmed spent a substantial part of the meeting explaining why the model-evaluation work should be visible in the final engineering story. The project does not have a large external user study, thousands of testers or a survey dataset from which to make strong empirical claims about user outcomes. The evaluation work gives the team a repeatable way to compare candidate model behavior against controlled cases and documented criteria instead of selecting a provider only because it was familiar or happened to be configured first.

The team saw this as useful academically because it turns the model choice into an engineering decision that can be explained and defended. It also helps answer a likely question in the presentation: why this production model, and how did the team decide it was suitable for the application?

7.2 What the evaluation can support

- Run the same representative test cases against candidate models.
- Compare behavior using agreed criteria rather than anecdotal impressions.
- Add more fixtures and statistics over time as the product evolves.
- Use the results to justify the release model while preserving the provider-independent architecture.
- Retain the evaluation framework as a foundation for future model changes rather than treating model selection as a one-time guess.

The discussion also explored a longer-term possibility in which TrustAI could choose among multiple candidate models based on comparative performance. That remained a future-product idea, not a claim about the current release.

7.3 No guarantee of model correctness

A team member challenged the word “guarantee,” noting that model-based systems can still produce incorrect or unsupported output. The group agreed that the product should be framed as best-effort decision support rather than a guarantee of correctness. This matches the safety and uncertainty language the team has already built into the product.

Presentation wording to preserve

The evaluation can justify why a model was selected and can show measured behavior on defined test cases. It should not be presented as proof that the system can never hallucinate or that every analysis is guaranteed to be correct.

8. Product Differentiation and Future Direction

The meeting included a candid discussion about the difference between the current MVP and what TrustAI could become. The team acknowledged that a single model response alone is not enough to claim a durable

product advantage. The value comes from the surrounding system: structured input, controlled analysis, validation, deterministic scoring, user history, Visual Inspection, explainable output, testing, provider abstraction and a deployment that packages those pieces into one usable buyer workflow.

The model-evaluation framework was discussed as one possible way to deepen that differentiation later. A future version could expand comparative testing, incorporate stronger grounding, or select models differently for different tasks. The group also briefly discussed the possibility of continuing the product after graduation if there is enough interest. No commercial commitment was made during the meeting.

9. Sprint Demonstration and Agreement Evidence

The team then turned to missing submission evidence. Ahmed Al-Mandalawi asked where the final signed Group Project Agreement and any original sprint-demo recordings were stored. The group did not establish during the meeting that a complete set of contemporaneous TrustAI sprint-demo recordings was available. A shared OneDrive location was recalled as a possible location for team documents, including the agreement, and the team agreed that the original artifacts should be located and verified before being treated as submission evidence.

The group also discussed whether a missing sprint demonstration could be recreated by checking out an older commit and showing the system as it existed at that point. This was discussed only as a possible explanatory reconstruction; it was not established as contemporaneous sprint-end evidence.

Evidence-handling rule for final documentation

Existing authentic sprint recordings should be located and indexed first. If the team creates any later reconstruction for explanation or archival purposes, it should be labeled honestly as a reconstruction of a historical commit rather than presented as contemporaneous sprint-review evidence.

This distinction matters because the Handbook expects a demonstration at the end of each sprint. The team's final documentation should preserve what actually happened rather than backfilling historical evidence without disclosure.

10. Trello Board Reconciliation

Trello cleanup was identified as one of the largest remaining submission tasks. Ahmed noted that some cards had no owner, some completed work was still sitting in active columns, and some work completed outside the original ticket structure was not represented on the board at all.

The team agreed that this cleanup cannot be done accurately by one person guessing who did what. Each member needs to review their own work and correct the board.

- Assign the correct owner to cards that were actually worked on.
- Move completed work to Done.
- Leave genuinely incomplete work in an appropriate open state.
- Create or update cards for material work that occurred outside an earlier ticket when that improves traceability.
- Use a Cancelled, Deferred, Future Work or similar state for work intentionally removed from the Capstone release.
- Keep improvement ideas separate from defects in the delivered MVP.
- Make sure the final board tells the same story as the repository and final documentation.

The group also discussed grader access. The board still needed a final access check so that the grader could review the live project-management record rather than rely only on screenshots. The exact access method remained to be confirmed and configured.

11. Final Testing, Defect Triage and Scope Discipline

A question was raised about what the team should do if UX, UI or functional testing discovers additional issues during the final days. The answer was effort-based rather than absolute. The team should create an issue, review its impact and decide whether it is worth changing the frozen release.

- Small, high-impact and presentation-blocking defects may still justify a fix.
- Large changes this late should generally not displace required submission work.
- Known non-blocking limitations can remain documented as future work.
- The presentation should focus on the agreed scope and what the team can demonstrate reliably.
- The written record should remain candid about what is complete, what is deferred and what is still open.

Part of the discussion was about not allowing secondary shortcomings to dominate the presentation. The practical resolution was not to hide defects; it was to distinguish the delivered scope from future enhancement work and to avoid spending the limited presentation time on issues that are outside the assessed user stories.

12. Academic Integrity and Source Attribution

The group briefly discussed source attribution and the need to review the final submission against Quantic's academic integrity requirements. The team agreed that any external material, code or other work that requires attribution should be handled in the final written documentation rather than improvised during the short presentation.

The exact form of that attribution was not settled during the call. The team agreed to rely on the Handbook and final submission requirements rather than guessing.

13. Final Presentation Format

13.1 One coordinated group recording

The team confirmed that the final submission should be one coordinated group recording, not five unrelated individual videos. The preferred approach is for all five members to be present on the same call, with the presentation passed from one person to the next.

The group considered whether individually recorded segments could be stitched together. The main concern was that the Handbook requires all members to be present and visible, and the team felt a heavily edited compilation could look less natural and make the ID/attendance requirement harder to satisfy. The safer plan was therefore to use individual recordings only as rehearsal material and complete the final presentation live in one coordinated session.

13.2 Time limit and speaking allocation

The team recalled that the presentation was approximately 15–20 minutes and agreed that the exact limit had to be checked against the rubric. With five members, the available time is tight. A rough working assumption was that most members would have about two to three minutes, with a little more time reserved for the introduction, transitions and close.

The Handbook confirms the 15–20 minute requirement. That means a five-minute segment for every member would be too long before the introduction and demonstration are even included.

13.3 IDs and camera presence

The group explicitly remembered that government-issued identification must be shown to the camera and discussed rehearsing that part as well. The intention was to remove uncertainty before the actual recording so the final session could run continuously within the allotted time.

14. Proposed Presentation Skeleton

The team did not lock a final speaker order during the meeting, but it agreed on the need for a single skeleton that prevents overlap and gives every member a distinct contribution. One suggested flow was:

1. Introduction: introduce the team, the marketplace-buyer problem and the TrustAI Marketplace solution.
2. Product / live user flow: show the application and the core user stories working in production.
3. Frontend: explain the main user experience, responsive design, results presentation and Visual Inspection interaction.
4. Backend / architecture: explain the FastAPI services, PostgreSQL persistence, provider abstraction, deterministic scoring and key design decisions.
5. Model evaluation / analysis quality: explain how candidate models were compared, why the production choice was made and how uncertainty is controlled.
6. Repository / CI/CD / deployment: show the engineering workflow, review controls, automated tests and AWS deployment pipeline.
7. Agile delivery / Trello: show how work was split, tracked, completed, deferred and reconciled across the sprints.
8. Close: summarize what was delivered, the strongest engineering lessons and clearly bounded future enhancements.

Because the total recording must remain between 15 and 20 minutes, not every item above can receive equal time. Ahmed volunteered to prepare the initial skeleton so the group could decide what to keep, combine or shorten before rehearsal.

15. Rehearsal Strategy

The meeting settled on a two-stage preparation approach. First, the team would prepare the skeleton and each member would draft or rehearse their own short section. Those individual practice recordings could then be reviewed before the team rehearsal. Second, the entire group would run the presentation together as one session to test timing, transitions, screen sharing and ID presentation.

- Skeleton first, so everyone knows what they are responsible for.
- Individual practice is allowed and encouraged for preparation.
- Use the practice recordings to tighten wording and remove repeated material.
- Run a full group rehearsal with the same sequence planned for the final recording.
- Time the whole presentation from beginning to end.
- Prepare all browser tabs, diagrams, repository views and Trello pages before starting.
- Use the final rehearsal to decide whether another polish session is needed before the submission recording.

16. Rehearsal Schedule

The largest scheduling discussion concerned the next team session. Time-zone differences again made a normal meeting slot difficult. Adrian had a Sunday commitment beginning around 1:30 PM local time, while Ahmed could accommodate a late Saturday evening in California more easily than the very early Sunday-morning hours that some CAT options would create.

After working through the options, the group settled on 7:00 AM CAT on Sunday for the team, which corresponds to approximately 10:00 PM PDT on Saturday for Ahmed. The team planned to pass the time to Samar and confirm any final availability through WhatsApp.

Next rehearsal target

Sunday, September 6, 2026 at 7:00 AM CAT / Saturday, September 5, 2026 at approximately 10:00 PM PDT.

17. Decisions and Working Agreements

ID	Decision / working agreement	Why it mattered
S03-01	Treat the application as functionally ready for Capstone submission unless serious new defects are found.	Protects the remaining time for evidence, documentation and presentation quality.
S03-02	Stop discretionary feature development for the submission release.	Prevents late scope from destabilizing the production build.
S03-03	Use the model-evaluation work as part of the engineering explanation for model selection.	Shows a repeatable decision process rather than an arbitrary provider choice.
S03-04	Do not describe model-based analysis as guaranteed correct.	The product remains decision support with explicit uncertainty.
S03-05	Reconcile Trello member-by-member against actual work completed.	One person cannot accurately reconstruct ownership and status for the entire team.
S03-06	Locate authentic sprint recordings and the signed agreement before attempting any reconstruction.	Preserves the integrity of the historical project record.
S03-07	Any later reconstruction of an old sprint state must be labeled as a reconstruction, not passed off as an original sprint-end recording.	Separates historical evidence from explanatory reproduction.
S03-08	Use one coordinated 15–20 minute final presentation with all members present, visible and speaking.	Matches the required group demonstration format.
S03-09	Use individual recordings only as rehearsal material unless the team later verifies a compliant editing approach.	Keeps the final recording simple and avoids camera/ID/continuity problems.
S03-10	Ahmed will prepare the initial presentation skeleton.	Gives the team a concrete structure for assigning non-overlapping sections.
S03-11	Hold the main rehearsal Sunday at 7:00 AM CAT / Saturday 10:00 PM PDT.	Creates a firm preparation milestone before the final recording.
S03-12	Triage newly discovered issues by impact and effort; large non-blocking work becomes future work.	Balances release quality with submission deadlines.

18. Action Items

Owner	Action	Timing	Status
Ahmed Al-Mandalawi	Prepare the first presentation skeleton and proposed flow, including a clear place for the model-evaluation story.	By the next day / before team rehearsal	Open
Ahmed Al-Mandalawi	Continue the rubric-to-repository documentation review and identify the highest-priority missing evidence.	Before rehearsal	In progress
Ahmed Al-Mandalawi	Maintain the September 3 meeting record as part of the final Capstone documentation set.	Immediate	In progress
All members	Test the current production application and raise only material defects that could affect submission or the demo.	Before rehearsal	Open
All members	Review their own Trello cards: owner, status, Done/Cancelled/Deferred/Future Work, and missing material work.	Before rehearsal	Open
Mulima Chibuye + team	Confirm the Trello access method required for the grader and make the final board accessible.	Before submission	Open
Team	Locate the signed Group Project Agreement in the shared storage; if not found, complete the required final page again with accurate details.	Before submission	Open
Team	Determine whether any original TrustAI sprint-demo recordings exist, verify any artifacts found, and index only evidence that can be attributed to the Capstone accurately.	Before submission	Open
Each presenter	Prepare a concise personal section based on the agreed skeleton and keep it within the allocated time.	Before Sunday rehearsal	Open
Entire team	Run the coordinated rehearsal, including screen share, transitions, camera visibility and ID presentation.	Sunday 7:00 AM CAT / Saturday 10:00 PM PDT	Scheduled
Entire team	After rehearsal, decide whether one additional polish session is needed before the final recording.	After rehearsal	Open

19. Risks and Final-Stage Concerns

Risk / issue	Concern	Response
Presentation overload	Five people, a live demo and the engineering story must fit in 15–20 minutes.	Use one skeleton, strict section ownership and timed rehearsals.
Trello drift	Missing owners, stale statuses and unrecorded work could weaken traceability.	Each member reconciles their own cards; final board is checked against repo/history.
Missing sprint evidence	Some recordings are local or not yet centrally indexed.	Locate authentic recordings first; do not mislabel reconstructions.
Agreement artifact	The team was unsure where the signed final agreement page was stored.	Search shared storage now; re-complete only if genuinely missing.
Late defects	Final testing can still expose issues days before recording.	Fix only serious/high-value issues; defer large non-blocking work.
Model overclaiming	Evaluation results could be misread as a guarantee of correctness.	Frame evaluation as comparative evidence and retain uncertainty/best-effort language.
Provider vs architecture confusion	A specific production model can look inconsistent with provider-agnostic design.	Explain production selection separately from provider abstraction.
Uneven presentation familiarity	Members may know their own area but not the full product story.	Use the group rehearsal to make all members comfortable with the whole flow.
Submission-access failure	Repository or Trello can be technically complete but inaccessible to the grader.	Perform a final access audit before submission.

20. Open Questions to Close Before Final Recording

- Where is the signed final page of the Group Project Agreement currently stored?
- Which original sprint-demo recordings exist, what sprint does each one correspond to, and where will they be indexed?
- Has the Trello board been made accessible to the grader and fully reconciled with Done, Cancelled/deferred and future work?
- What exact five-person speaker order will be used in the final presentation?
- How much time will be allocated to the live application demonstration versus architecture, testing, model evaluation, CI/CD and project management?
- Which exact production cases will be demonstrated so the team shows a range of inputs without risking an unpredictable live result?
- Which diagrams and repository views are essential enough to show within the 15–20 minute limit?
- Does the team need one additional rehearsal after Sunday before recording the final submission?

21. Next Meeting

The next session is the presentation walkthrough and rehearsal rather than another general development meeting. The group agreed on Sunday morning CAT time, with the final working time settled at 7:00 AM CAT, corresponding to late Saturday evening for Ahmed in California. Final attendance confirmation will be handled through the team's normal WhatsApp communication.

Planned agenda:

- Confirm that the production application remains stable.
- Review the presentation skeleton and final speaker order.

- Confirm each person's section and remove overlap.
- Review the model-evaluation explanation and keep it concise and defensible.
- Confirm Trello cleanup and remaining submission evidence.
- Run the complete presentation once without stopping.
- Practice camera/ID presentation and screen-share transitions.
- Review timing and decide what must be shortened, moved to documentation or removed.
- Set the final recording date and any last polish session.

22. Closing Summary

The September 3 meeting moved the team from final implementation and production checking into submission closeout and presentation rehearsal. The central conclusion was that TrustAI Marketplace did not need another round of broad feature development; the remaining effort should concentrate on validating, organizing and presenting the system already delivered.

Ahmed's production update gave the group confidence that the latest application, including the standard text-analysis flow and Visual Inspection, was ready for final team testing. The model-evaluation work was identified as an important piece of the engineering story because it provides a repeatable basis for explaining model selection while also making clear that no model can be treated as infallible.

The meeting also surfaced the less visible work that can make or break the submission: Trello accuracy, sprint-demo evidence, the signed agreement, the design-and-testing document, grader access and the discipline to stop changing a working release. The team agreed that those items now deserve the same attention that implementation received earlier in the project.

Finally, the presentation moved from a general idea to a concrete plan. Ahmed would prepare the initial skeleton; each member would prepare a short, distinct section; individual practice could be used to improve delivery; and the final submission would be rehearsed as one coordinated group presentation. The next milestone was set for Sunday at 7:00 AM CAT, giving the team a firm point at which the application, evidence and presentation story all needed to come together.